

Implementation and Performance of Barnes-Hut N-body algorithm on Extreme-scale Heterogeneous Many-core Architectures

Journal Title
XX(X):1–13
©The Author(s) 2016
Reprints and permission:
sagepub.co.uk/journalsPermissions.nav
DOI: 10.1177/ToBeAssigned
www.sagepub.com/

SAGE

Masaki Iwasawa^{1,7,2}, Daisuke Namekata², Ryo Sakamoto³, Takashi Nakamura⁴, Yasuyuki Kimura⁵, Keigo Nitadori², Long Wang^{6, 2}, Miyuki Tsubouchi², Jun Makino⁷, Zhao Liu^{8,9}, Haohuan Fu¹⁰, Guangwen Yang¹¹

Abstract

In this paper, we report the implementation and measured performance of our extreme-scale whole planetary ring simulation code on Sunway TaihuLight and two PEZY-SC2 systems: Shoubu System B and Gyoukou. The numerical algorithm is the parallel Barnes-Hut tree algorithm, which has been used in many large-scale astrophysical particle-based simulations. Our implementation is based on our FDPS framework. However, the extremely large numbers of cores of the systems used (10M on TaihuLight and 16M on Gyoukou) and their relatively poor memory and network bandwidth pose new challenges. We describe the new algorithms introduced to achieve high efficiency on machines with low memory bandwidth. The measured performance is 47.9, 10.6 PF, and 1.01PF on TaihuLight, Gyoukou and Shoubu System B (efficiency 40%, 23.5% and 35.5%). The current code is developed for the simulation of planetary rings, but most of the new algorithms are useful for other simulations, and are now available in the FDPS framework.

Keywords

HPC, N -body simulation, planetary ring, heterogeneous many-core architecture, Sunway TaihuLight, PEZY-SC

1 Introduction

The architecture of HPC platforms has shown significant changes in the last three decades, from vector-parallel architecture to distributed-memory scalar processors, and then many-cores with SIMD units, accelerators, and most recently heterogeneous many-cores.

The ideas behind the accelerators and heterogeneous many-cores are very similar. In both cases, a number of relatively simple and thus energy- and area-efficient processors are combined with relatively small number of complex, high-performance processors. The difference is that in the case of the accelerator systems, the complex, high-performance cores are usually commodity processors (the host CPU) in one die, and accelerator cores are in a separate die, and they are connected by general-purpose link such as the PCI express. The host CPU is usually of x86 architecture. GPGPUs used with x86 processors are the currently the most widely used accelerator systems. This architecture has many advantages, but the most important one is that all of the hardware and software for the host CPU is already there. The developer of the accelerator hardware can concentrate on the accelerator hardware itself and software to make use of it.

The disadvantage of the accelerator system is the existence of the communication link between the host and accelerators. Usually, both the bandwidth and latency of the communication between the host and accelerators is limited by that of the standard PCIe specification. The theoretical peak throughput of the PCIe 3.0 standard (with 16 lanes) is only 16GB/s. On the other hand, some of latest GPGPUs have multiple channels of the HBM memory, with

the total bandwidth approaching to 1TB/s. If we divide the calculations of an application into that on the host and that on accelerators, communication between them would become necessary, and in many cases that would limit the performance. Thus, either we have to port all of application to the accelerator side, or we have to live with relatively low performance. In addition, in many cases size of the on-board memories of the accelerator boards is small, less than 16GB, and thus either the problem size is limited or we need to store the data in the main memory of the host CPU.

In principle, heterogeneous many-cores can solve these limitations of accelerator systems, since what correspond to the host CPUs of accelerator systems are now integrated to the same LSI chip as the accelerators, and they share

¹National Institute of Technology, Matsue College

²RIKEN Center for Computational Science

³PEZY Computing K. K.

⁴Preferred Networks Inc.

⁵ExaScaler Inc.

⁶Department of Astronomy, School of Science, The University of Tokyo

⁷Graduate School of Science, Kobe University

⁸Department of Computer Science and Technology, Tsinghua University

⁹National Supercomputing Center in Wuxi

¹⁰Department of Earth System Science, Tsinghua University

¹¹Department of Computer Science and Technology, Tsinghua University

Corresponding author:

Masaki Iwasawa, National Institute of Technology, Matsue College, 14-4 Nishiikuma-cho, Matsue, Shimane, 650-0047, Japan

Email: masaki@matsue-ct.jp

the same physical memory. Examples of such architecture include Sunway SW26010 and PEZY-SC2. The former is used in the TaihuLight system, which was ranked #1 in the Top 500 list four times in years 2016 and 2017. The latter was used in the Gyoukou system, which was ranked #4 in the Top 500 list of November 2017. We can see that though the heterogeneous many-core architecture has clear advantages, they are yet to be widely used.

Currently, to port applications to these systems and achieve good performance requires quite a lot of efforts. One reason is that “automatic” parallelization through, for example, OpenACC results in rather poor performances in the case of TaihuLight (Cai et al. 2018), and currently only a subset of OpenCL is available on PEZY-SC2. In the case of PEZY-SC2, right now the host CPU on chip is disabled and thus the current system is an accelerator system with a Xeon-D CPU. However, since the memory on PEZY-SC2 is larger than that on Xeon-D, all data can (and should) be on the side of PEZY-SC2.

Another reason is that the performance ratio between the complex cores and simple cores tend to be very large. In the case of TaihuLight, one complex core and one simple core have the same peak performance, but there is only one complex core for every 64 simple cores. Thus, even though it is not easy to write codes for simple cores, almost all codes should be moved to the simple core side to obtain decent performance.

One way to make the application development on these machines easier is to provide DSLs or frameworks, in which the users express the problem to be solved or numerical method to be used in high-level, machine-independent way. The machine-specific part of DSL runtime library must be optimized to each architecture, but since one DSL can be used to implement many applications, if such a DSL is possible a lot of works by many researchers can be used for more productive researches.

We have been developing FDPS, Framework for Developing particle simulators (Iwasawa et al. 2016). The basic idea of FDPS is to provide a set of library functions necessary for high-performance, highly scalable particle simulation codes. FDPS receives from user applications the definition of particle data (originally in C++ class) and a function to evaluate particle-particle interactions as the source code. FDPS itself is written as a template library which is compiled with these user-defined class and interaction function. Thus, a user’s application can call functions defined in FDPS to process particles they defined and to calculate the interactions they defined. We have extended the API so that FDPS can accept the particle data class (or struct) and the interaction function written in both Fortran (Namekata et al. 2018) and C. The pure-C language interface makes it possible for programs written in any language with reasonable FFI to C language to use FDPS functions.

FDPS relies on parallel Barnes-Hut tree algorithm with domain decomposition by multisection algorithm and local essential tree method (Makino 2004) for interaction calculation. Currently, FDPS supports usual multicore architectures and also accelerator architectures. In the case of the support of accelerator architectures, the multiwalk algorithm (Hamada et al. 2009) is used. This means that

everything other than the interaction calculation using the interaction list is done on the host side, and the user-supplied interaction function need to take care of the data transfer between the host and accelerators.

Thus, currently FDPS does not support heterogeneous many-cores very well. One could use the accelerator support, but the performance gain would be rather limited because of the reasons described above. In order to improve the efficiency, it is necessary to move operations other than the interaction calculation such as the tree construction and the construction of the interaction list to the simple cores.

In this paper, we report the result of porting relatively simple N-body simulation code for planetary ring systems, developed based on our FDPS framework, to two heterogeneous many-core processors: Sunway SW26010 and PEZY-SC2. The code is not yet the complete port of FDPS to these processors, but more like a production code based on FDPS for a specific problem. The reason why full FDPS port is not yet done is simply that we made this porting partly to evaluate the architecture and partly to try large-scale calculations which were not practical on other architectures.

In the rest of this paper, we first give an overview of the current state of the arts for the large-scale simulations of planetary rings in section 2, and then short description of the architecture of the two systems in section 3. In section 4, we describe in detail the new algorithms we developed to achieve high performance on extreme-scale heterogeneous many-core architectures. In section 5 we describe how the performance was measured and achieved performance. Section 6 is for discussion and summary.

2 Simulation of planetary rings

Saturn’s ring was first observed by Galileo Galilei in 1610. For more than three centuries, it had been the only known ring system within our solar system. In 1977, rings of Uranus were found through occultation observations from an aircraft, and then in 1979 rings of Jupiter by Voyager 1 and in 1989 those of Neptune by Voyager 2. Very recently, it turned out that some of minor planets also have rings. The first distinctive example is 10199 Chariklo, which is one of Centaurs*. There are probably more Centaurs with rings. Thus, quite recently, a wide variety of ring systems have been found. How these rings were formed and have evolved is an important question in planetary science, and large-scale, global simulation, if possible, would help greatly to advance our understanding.

Planetary rings are usually at the radii around the Roche limit. Thus, mutual gravity between particles does not easily lead to the formation of new satellites, but is important enough to form spiral waves (“wakes”) in very small scales, which increase the effective viscosity and should enhance the radial transport of the angular momentum. On the other hand, the actual ring system seems to consist of very large number of narrow rings, separated with distinct gaps. It is believed that these gaps are maintained by high-order resonances with small embedded satellites (so-called moonlets), but whether or not such gaps can be formed and maintained by resonances has not been fully understood.

*Minor planets whose orbit is between those of Saturn and Uranus

Up to now, most of simulations of ring structures have been local ones, in which a small patch was cut out from the ring and simulated under the assumption of the local Hill approximation and periodic boundary condition (Wisdom and Tremaine 1988). Rein and Latter (Rein and Latter 2013) performed “Large-scale” simulation of viscous overstability in Saturn’s rings, using up to 204,178 particles and up to 10,000 orbits using this local approach. Because very long simulations are necessary, the number of particles has been small. They used REBOUND (Rein and Liu 2012), an MPI-parallel N -body simulation code.

Michikoshi and Kokubo (Michikoshi and Kokubo 2017) performed global simulations of rings with the largest number of particles reported so far. They used 300M particles to model two narrow rings of Chariklo. They have developed their parallel code using the framework we developed, FDPS (Iwasawa et al. 2016; Namekata et al. 2018). They so far followed the system only for 10 orbital periods.

The total calculation cost is roughly proportional to number of particles multiplied by the number of orbital periods followed, since the calculation cost per timestep is $O(N \log N)$ when Barnes and Hut tree algorithm is used and the number of timestep required for ring simulations is essentially independent of the number of particles. Thus, we can conclude that the size of state-of-the-art simulations of planetary rings is around 10^9 particle-orbits, or around 10^{12} particle-steps because 10^3 steps are necessary for one orbit.

We should note that even though the simulations so far done in this field is relatively small, that does not mean there is no need or possibilities for larger scale simulations. If we want to model the global structures of rings, we cannot rely on local treatment. For example, the effect of resonances with small satellites can only be studied using global simulations. On the other hand, the number of particles one need for global simulations, even for a very narrow radial range, is very large. For example, consider the A ring of Saturn with the radius of around 1.3×10^5 km. The typical radius of ring particles is 6 m (Zebker et al. 1985), and the optical depth of the ring is around unity. Thus, we need 10^4 particles per square km or around 10^{12} particles for the radial range of 100 km. With this radial range, we can model many of fine features observed by Cassini directly.

If we could use particles with larger size, we could reduce the number of particles required significantly. However, that would change the viscous diffusion timescale of the ring, and thus what would be observed. It is necessary to perform simulations with particles of real physical radius, which would require at least 10^{16} and ideally 10^{19} particle steps.

In other fields of astrophysics, very large simulations have been performed. For example, Ishiyama (Ishiyama 2014) used 4096^3 particles to follow the formation and growth of dark matter halos of smallest scales. This simulation corresponds to 10^{16} particle steps. Part of this calculation was performed on K computer. The performance of K computer is 4.0×10^{10} particle steps per second on the entire K computer, or 60,000 particle step per second per core for a processor core with the theoretical peak performance of 16 Gflops (Ishiyama et al. 2012). The efficiency they achieved is 55% of the theoretical peak.

The algorithms used in large-scale N -body simulations are rather similar, and that means they are well studied and

close to optimal. All of them use domain decomposition and Barnes and Hut tree algorithm. For domain decomposition, several variations have been used, such as Orthogonal Recursive Bisection (Salmon et al. 1990), Hashed Oct Tree (Warren and Salmon 1992), Multisection (Makino 2004).

Efficient implementations on large-scale GPGPU clusters exist (Hamada et al. 2009; Portegies Zwart and Bédorf 2014; Bédorf et al. 2014). Bédorf et al. (2014) performed the simulation of Milky Way Galaxy using 2.42×10^{11} particles. The achieved performance is 24.77 PF on ORNL Titan, and one timestep took 5.5 seconds. Thus they have achieved the performance of 4.4×10^{10} particle steps per seconds. The theoretical peak performance of Titan is 73.2 PF in single precision. Thus, the achieved efficiency is 33.8%.

3 Sunway TaihuLight and PEZY-SC2 systems

In this section, we briefly describe the features of the Sunway TaihuLight system and two systems with PEZY-SC2 processors: Gyoukou and Shoubu System B. In table 1, we summarized the features of the systems. For more details of TaihuLight system see Fu et al. (2016). TaihuLight consists of 40960 Sunway 26010 processors, and Gyoukou and Shoubu System B 13312 and 512 PEZY-SC2 processors, respectively. Unfortunately, Gyoukou was turned off by March 31, 2018, and thus our performance measurement on Gyoukou system was based on a preliminary version of the simulation code, and the efficiency measured on Gyoukou is lower than that measured on Shoubu System B.

One SW26010 processor consists of four CGs (core groups), each with one MPE (management processing element) and 64 CPEs (computing processing elements). Both MPE and CPE are 64-bit RISC cores. MPE has L1 cache memories for both instructions and data, and also L2 data cache. On the other hand, each CPE has L1 instruction cache and 64KB of local data memory. CPEs can communicate with the main memory through DMA. Each CPE can initiate multiple asynchronous DMA operations.

Each core group is connected to 8GB DDR3 DRAM memory with the theoretical peak transfer rate of 34GB/s. The processor runs at the clock frequency of 1.45GHz, and each core (both MPE and CPE) can perform four double precision FMA operations. Thus, the theoretical peak performance of one processor is 3016 Gflops and that of one CG is 754 Gflops. Thus, even when we use the nominal number for DRAM bandwidth, the B/F ratio is only 0.045. This is less than 1/10 of the number for K computer.

This memory B/F value is very small because in the original tree algorithm, the required B/F is more than unity. By using the algorithm developed by Barnes (1990), the required B/F is reduced. For example, if the interaction list is shared by a few hundred particles, the dominant procedure of memory access is the tree construction, in which each particle data is read or written about 50 times per step. The typical size of a particle is about 64B, therefor the amount of memory access per step per particle is roughly 3200. For the force calculation in the Saturn’s ring simulation, the size of the interaction list is about 2000 and the number of operations per interaction is about 40. Thus

the total number of floating operations per step per particle is 80000. Considering random memory access for the tree construction, the required B/F should be much greater than 0.04.

Compared to that of K computer, the network is also weak, with the total bandwidth of around 10GB/s per node. This is about the same as the performance of a single link of 6D torus network of K computer. Since the SW processor is around 25 times faster than the SPARC64 processor of K computer, the relative network bandwidth is different by more than two orders of magnitudes.

One PEZY-SC2 processor chip consists of 2048 processors (64 of them are disabled and the available number of processors is 1984). Each processor can perform 1, 2 and 4 multiply-and-add operation for FP64, FP32, and FP16 data. For FP32 and FP16, 2- and 4-way SIMD operations are performed. With the clock speed of 700MHz, the theoretical peak speed is 2.8, 5.6 and 11.1TF, for FP64, FP32 and FP16, respectively. At present, each SC2 processor chip have 4 channels of DDR4 memory, for the peak throughput of 76.8GB/s. Thus B/F is 0.027.

They have three levels of shared cache, but without coherency. Instead, they have explicit cache flush instructions to each levels. Two processors share L1D, and 16 processors L2D, and all processors LLC. Each processor runs either four or eight threads simultaneously. Thus, it is relatively easy to hide the latency of the arithmetic units and L1D.

In the original design, each SC2 processor chip had six MIPS64 cores, which were supposed to run the operating system and main body of the application programs. Unfortunately, currently they are disabled, and operating system and application programs run on the frontend Xeon D-1571 processor. Each Xeon D hosts eight SC2 processors. Thus, the performance ratio between Xeon D and SC2 is close to 100. Moreover, these eight SC2 are connected to Xeon D through single PCIe Gen3 16-lane channel. Thus, the peak data transfer speed between one SC2 and Xeon D is 2 GB/s, for the peak speed of 2.8TF.

In summary, TaihuLight and two systems based on PEZY-SC2 processors share the following characteristics:

1. Very large performance ratio between “general-purpose” and “computing” cores, close to 1:100.
2. Very small memory B/F numbers, around 0.03.
3. Even smaller network B/F numbers, 0.003 or 0.0006.
4. Very large number of MPI processes, 160k or 10k.
5. Very large number of “computing” cores per MPI process, 64 or 1984.

Just one of these characteristics makes it very difficult to achieve reasonable performance for particle-based simulations using previously known parallelization algorithms. In the next section, we describe the new algorithms we implemented to achieve good performance on these systems.

4 New algorithms for extreme-scale simulations

4.1 Overview of new algorithms

In this section, we describe the new algorithms we made in order to utilize TaihuLight and PEZY-SC2 based systems

for the simulations of self-gravitating planetary rings. The following is the list of new algorithms.

1. The reuse of the interaction list over multiple timesteps.
2. Elimination of the global all-to-all communication.
3. “Semi-dynamic” load balance between computing cores
4. Optimizations specific to the ring geometry.

For TaihuLight and PEZY-SC2 based systems, we have modified our FDPS framework in architecture-specific way so that we implement the algorithms and run the code under the limited available time and software environment. However, many of these algorithms are ported back to the original FDPS so that anybody who uses FDPS can take advantage of these new algorithms.

In the rest of this section, we briefly describe these new innovations.

4.2 Reuse of the interaction list

The following gives the usual steps for highly parallel code for self-gravitating particle system:

1. Perform domain decomposition.
2. Exchange particles so that particles belong to appropriate domains.
3. Perform interaction calculation using fast algorithm such as Barnes-Hut tree.
4. Integrate the orbits of particles.
5. Go back to step 1.

In the case of approaches with local essential tree, step (3) consists of the following substeps:

- (3a) Construct the “local” tree structure from particles in the domain.
- (3b) Collect the information necessary for the calculation of interaction from other processes (so called local essential tree).
- (3c) Construct the “global” tree from the collected information.
- (3d) For small groups of particles, traverse the tree and calculate the interaction. Repeat this for all groups.

In the original algorithm (Barnes and Hut 1986), the traversal of the tree is done for each particle, and force calculation is done during the traversal. However, on almost all modern implementation (Ishiyama et al. 2012; Bédorf et al. 2014), following the idea of Barnes (Barnes 1990), tree traversal are done for groups of neighboring particles, which are constructed using the tree structure itself. During the traversal for a group, the list of particles and tree nodes which exert the force on this group of particles is constructed, and actual force calculation is done through the double loop over particles in the group and those in the list. This structure makes it possible to use vector pipelines, scalar SIMD units, and even special-purpose computers (Makino 1991) with high efficiency. For GPGPUs, the extension of

System	Processor	# of MPI processes	# of cores per process	memory B/F	network B/F
TaihuLight	SW26010	160k	64	0.045	0.003
Gyokou	PEZY-SC2	8192	1984	0.03	0.0006
Shoubu System B	PEZY-SC2	512	1984	0.03	0.0006

Table 1. Parameters of HPC systems we used.

this algorithm, in which multiple lists are constructed and then sent to GPGPU, is used (Hamada et al. 2009).

This approach does not work well on TaihuLight or PEZY-SC2 based systems, because of the low performance of general-purpose core and limited memory bandwidth. The performance we can achieve with either approach for ring simulation on these machines is less than 1%. Thus, it is necessary to reduce the cost of tree construction and tree traversal, and we achieved this by using the same interaction lists over multiple timesteps. We call this method the persistent interaction list method.

The idea behind this method is essentially the same as that for the neighbor-list method used in many simulation codes for particles with short-range interactions. In the case of the ring simulation, in which the dominant motion of particles is the differential rotation, the persistent interaction method gives additional force error from the deformation of the tree cells. We will discuss this error and a solution to reduce this error with less additional cost in section 6.1.

By using this persistent interaction list, we can reduce the calculation cost of the part other than the interaction calculation drastically. While we are using the same interaction lists, we skip the domain decomposition, exchange of particles, construction of the local tree. We still need to update the physical quantities of the nodes of the tree, since particles move at each timestep. We first update the information of the nodes of the local tree. Then, using the list of nodes for the local essential tree, the communication between the nodes is performed. Finally, the physical quantities of the global tree are updated, and the force calculation is performed using this updated global tree and the persistent interaction list.

The most time-consuming part of the tree construction is the sorting. In the case of TaihuLight, we implemented the parallel sample sort (Blueloch et al. 1991) on CPEs. In the case of Gyokou, the sorting was performed on Xeon D host processor. In the case of Shoubu System B, it was performed on the side of PEZY-SC2 processors. Also, some other operations are moved from Xeon-D to PEZY-SC2. Thus, the overall performance is significantly better for Shoubu System B. As stated earlier, Gyokou was turned off on March 31, 2018, and we could not measure the performance of our improved code on systems with more than 512 PEZY-SC2 processors.

We have ported all operations in timesteps in which the interaction list is used (list-reusing steps), except for MPI functions for communication, to CPEs (TaihuLight) or SC2 processors (PEZY-SC2 based systems). For the timestep in which the interaction list is constructed (list-constructing step), some of operations are still done on Xeon D in the case of PEZY-SC2 based systems.




Figure 1. Schematic figure of domain decomposition by the multisection method in x - y coordinate. Domains are divided by 16×16 .




Figure 2. Schematic figure of domain decomposition by the multisection method in cylindrical coordinate. Domains are divided by 2×128 .

4.3 Tree and Domain structures on Cylindrical Coordinate

We want to model a relatively narrow ring, and this means the usual domain decomposition in Cartesian coordinates can cause serious problems. Figure 1 illustrates the problem. We can see the domains near the y axis are very elongated. This irregular shape of domains results in the increase of communication between processes, and thus serious degradation in the efficiency.

We can avoid this problem, if we apply the domain decomposition in the cylindrical coordinates (figure 2). The amount of data size of communication is roughly proportional to the surface areas of domains. For example, if we use 160K (5×32000) nodes, the total surface area of domains in Cartesian coordinate is about 20 times larger

than that in Cylindrical coordinate. Thus the amount of data size of the communication is reduced by a factor of 20. Note that we can also use the cylindrical coordinates for the construction of the tree. Since the ring is narrow, the local distance s in the Cartesian coordinates (x, y, z) can be approximated by that in cylindrical coordinates (r, ϕ, z) .

$$ds^2 = dx^2 + dy^2 + dz^2 \sim d\phi^2 + dr^2 + dz^2, \quad (1)$$

when $r \sim 1$. Thus, we can use the cylindrical coordinate for domain decomposition and tree construction and even for the tree traversal, without any modification of the algorithm or program itself. The actual interaction calculation is faster in Cartesian coordinates and thus Cartesian coordinates is used.

4.4 Coordinate rotation

The simulation of ring with very large number of processes poses new challenges. As we increase the number of processes, the size of the domains becomes smaller. On the other hand, the timestep does not become much smaller even when we increase the total number of particles, since the random velocities of ring particles become smaller when we increase the number of particles. Thus, the distance that particles move can be comparable or even larger than the domain size, resulting in the increase in the amount of communication.

We can “solve” this problem by the rotation of the coordinates and domain structure, so that particles do not move much. If we rotate the coordinates at the speed of Kepler rotation at the center of the ring, particles at the center of the ring do not move much. Particles at other radial positions still move, but the speed becomes much smaller than that of the Kepler rotation. For example, when we use the same interaction list for 64 steps, each particle rotates for about 0.4 radian in the theta direction during the 64 steps. If we use 160K (5x32000) nodes, without coordinate rotation, all particles in all nodes need to be sent to other nodes. However, in the rotating frame, particles rotate very slowly. Even the fastest particles rotate by only 2×10^{-4} radian in the theta direction during 64 steps. As a result, by using coordinate rotation method, only 5% of particles need to be sent to other domains. Thus we can reduce communication cost by a factor of 20. Thus, communication due to Kepler rotation can be almost eliminated.

4.5 Elimination of all-to-all communication

In FDPS, the exchange of LET (local essential tree) data is done though a single call to the `MPI_Alltoallv` function. This implementation works fine even for full-node runs on K computer, but becomes problematic on systems with relatively weak network like TaihuLight and PEZY-SC2 based systems. We can eliminate this all-to-all communication, by constructing the “tree of domains” locally and let only higher-level information be sent to distant processes.

In the current implementation specialized to narrow rings, we implemented a very simple two-level tree, in which the second-level tree nodes have all processes in the radial direction. For example, if we have a process grid of (1000, 10), where 1000 in angular and 10 in radial direction, 10 domains in the radial direction are combined to one tree

node, resulting in 1000 second-level nodes. Only these 1000 nodes exchange their center-of-mass information. All LET information other than these center-of-mass data of second-level nodes are sent either to other second-level nodes (and then broadcast to lower-level nodes) or sent directly to lower-level nodes.

In this implementation, there is still one global communication in the angular direction, but we can use `MPI_Allgather` since only the top-level data are sent. `MPI_Allgather` is much faster than `MPI_Alltoall` because the performance of the former is limited by the injection bandwidth and that of the latter is limited by the bisection bandwidth. For example, on TaihuLight, the time for `MPI_Alltoall` with a short message size among 10000 processes is can be as much as several minutes. On the other hand, that for `MPI_Allgather` is about ten millisecond. Thus the reduction rate is roughly 10000. Thus the reduction in the communication was quite significant.

4.6 Load Balance among computing cores

In our current implementation, interaction lists are created at the list-construction step, and are reused for several steps. The total number of lists in one MPI process is around 10^5 , and we need to use 64 or 1984 computing cores efficiently for them. Therefore, some load balance strategy is necessary. We applied the following simple algorithm.

1. Sort the interaction lists by their length.
2. Assign the longest 64 lists on 64 CPEs (in case of TaihuLight).
3. For each remaining list, assign it to the CPE with the shortest total calculation cost.

Since the calculation time of cores is quite predictable, this algorithm works very well. Without this method, the typical variation of the force calculation cost among all CPEs is 40%. On the other hand, with this method, the variation is a few percent. As a result, by using this method, the time for the force calculation is reduced by 10%..

In the case of PEZY-SC2 based systems, we further improved the load balance by using multiple cores which share the cache for one interaction list.

4.7 Interaction Kernel

In the case of TaihuLight, we found the compiler-generated code for the interaction kernel, even when SIMD operations are used, does not give very good performance because there is no loop unrolling and the instruction scheduling is not ideal. As a results, with the compiler-generated code the achieved efficiency was around 10%. Thus we rewrite the interaction kernel fully in the assembly language, with hand-unroll and careful manual scheduling. As a result, we achieved more than 50% of the theoretical peak performance for the kernel.

We have applied similar optimization also on PEZY-SC2 based systems. In addition, on PEZY-SC2 based systems we used single-precision calculation for the interaction kernel. The force error introduce by using single-precision is much smaller than that by Barnes-Hut tree algorithm. For example, The relative force error by introducing center-of-mass approximation (effectively dipole moment

approximation) with tree opening-criterion $\theta = 0.5$ is about 0.1% (McMillan and Aarseth 1993). This force error can be expressed by 10 bits. Thus we can safely use single-precision calculation for the interaction kernel.

In order to avoid the large roundoff at the first subtraction of the position vectors, both positions and velocities are shifted with the new origin at the position of one of the particles which share the interaction list. After this shifting, positions and velocities are converted to single precision, and actual interaction calculation is done using single-precision SIMD operations.

5 Measured performance

5.1 How the performance is measured

To measure the performance, we measure the time for 64 timesteps, including the time for diagnostics. The execution time is measured by the MPI wallclock timer, and operation count is from the counted number of interactions calculated. Equation 2 gives the definition of the particle-particle interaction.

$$\mathbf{F}_{ij} = \begin{cases} -G \frac{m_i m_j}{r_{ij}^3} \mathbf{r}_{ij} & (r_{ij} > r_{\text{coll}}) \\ \left[-G \frac{m_i m_j}{r_{\text{coll}}^3} - \frac{m_i m_j}{m_i + m_j} \times \left(\kappa \frac{r_{ij} - r_{\text{coll}}}{r_{ij}} + \eta \frac{\mathbf{r}_{ij} \cdot \mathbf{v}_{ij}}{r_{ij}^2} \right) \right] \mathbf{r}_{ij} & (r_{ij} \leq r_{\text{coll}}) \end{cases} \quad (2)$$

with $\mathbf{r}_{ij} = \mathbf{r}_i - \mathbf{r}_j$, $\mathbf{v}_{ij} = \mathbf{v}_i - \mathbf{v}_j$, $r_{ij} = \|\mathbf{r}_{ij}\|$.

Here, \mathbf{F}_{ij} is the acceleration of particle i due to particle j , \mathbf{r}_{ij} and \mathbf{v}_{ij} are the relative position and velocity vectors, G is the gravitational constant (taken to be unity in this paper), m_i is the mass of particle i , r_{coll} is the distance at which two particles collide, and η and κ are parameters which determine the coefficient of restitution. We chose these parameters so that the coefficient of restitution in radial direction is 0.5.

We used this form to calculate all particle-particle interaction. For particle-tree-node interaction, we used center-of-mass approximation with $\theta = 0.5^\dagger$. Particle-particle interaction consists of 9 multiplications, 8 additions, and one square root and one division operations. Instruction set of Sunway 26010 processor does not include fast approximation for neither square root or reciprocal square root. So we implemented fast initial guess and high-order convergence iteration in software. The number of operations in this part is 7 multiplications, 5 additions and two integer operations. Therefore, for particle-cell interactions the number of floating-point operations is 31, and for particle-particle interactions, which include the repulsive force during physical collisions, is 49. The total number of floating-point operations is obtained by counting the number of interactions calculated and multiply them with these number of floating-point operations per interaction. We ignore all operations other than the interaction calculation, since as far as the number of floating-point operations is concerned, that for interaction calculation is more than 99% of total operation count.

For PEZY-SC2 based systems we used the same operation count as we used for TaihuLight, in order to make the

Table 2. Initial condition for weak scaling runs

Central planet	Saturn
Ring inner radius	10^5 km
Ring width	100 km
Number of MPI processes	1024 – 160,000
Number of particles per process	10^7
particle radius	3.5 – 500 m




Figure 3. Time per timestep for weak-scaling test. The number of particles per process is 10M. Triangles, crosses and squares show the results on TaihuLight, Gyoukou and Shoubu System B, respectively.

direct comparison possible, even though the details of the implementation of the force kernels are different[‡].

For the weak-scaling measurement, we have performed runs with 10M particles per MPI process on TaihuLight and PEZY-SC2 based systems. Initial condition is such that the ring width and ring radius is unchanged. Table 2 summarizes the initial condition.

5.2 Performance Results

Figures 3 and 4 shows the time per one timestep and the performance, for the weak scaling measurements. We can see that the weak scaling performance is quite good on both of TaihuLight and PEZY-SC2 based systems. The peak performance of single MPI process is roughly four times faster for PEZY-SC2 based systems, and that's the reason why they are three to four times faster in this weak-scaling measurement. We can see that Shoubu System B is about 50% faster than Gyoukou. As we have already discussed, this

[†]If we use quadrupole moment, we can use $\theta = 1.0$ to keep the same integration error as that with center-of-mass approximation with $\theta = 0.5$. In the case of $\theta = 1.0$, the number of interactions per particle is reduced by 30% compared with the case of $\theta = 0.5$ and the time to solution would be reduced by 30%. However, even if we use $\theta = 1.0$, the size of the interaction list is enough long so that the processors are always busy. Thus if we use quadrupole moment, the performance would not be changed.

[‡]PEZY-SC2 has special function units (SFU) for reciprocal square root operation and we used it for the force calculation. However, the latency of SFU is very large. If we considered this latency, the efficiency for PEZY-SC2 based system would increase by several percent.




Figure 4. Performance in petaflops for weak-scaling test. The number of particles per process is 10M. Triangles, crosses and squares show the results on TaihuLight, Gyoukou and Shoubu System B, respectively.

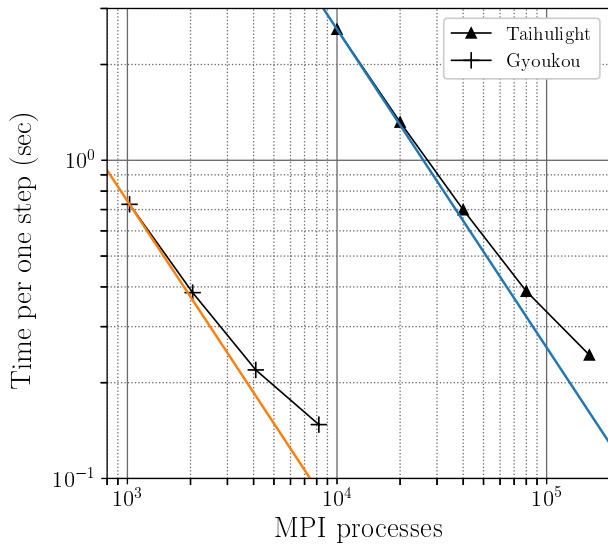


Figure 5. Time per timestep for strong-scaling test. Triangles and crosses show the results on TaihuLight and PEZY-SC2 based systems, respectively. The total number of particles is 10^{11} on TaihuLight and 10^{10} on PEZY-SC2 based systems. Solid lines indicate idealized scaling.

is not due to any hardware difference but the difference in the software used.

In figure 3, we can see that up-down features for the results on TaihuLight and Shoubu-B. This up-down is caused by the sudden change in the size of the particle group for Barnes' algorithm. Since the distribution of particles is almost uniform, all tree cells in the same level have nearly the same number of particles. When we apply Barnes' algorithm, we set the maximum number of particles in the group. The actual number is the maximum number of particles in tree cells, which does not exceed the limit. For example, when we set this limit to 1024, actual value can be anywhere between 1024 and 256. If it is, say, 300 for certain value of the total number of particles, when we double the number of particles, the actual value would become 600, and when we double the

Table 3. Breakdown of calculation time for weak-scaling runs

System	# processes	interaction	comm.	others
TaihuLight	10000	2.07	0.041	0.466
	20000	1.63	0.040	0.590
	40000	2.13	0.064	0.478
	80000	1.71	0.053	0.630
	160000	2.31	0.090	0.476
Gyoukou	1024	0.332	0.114	0.281
	2048	0.392	0.121	0.235
	4096	0.355	0.143	0.289
	8192	0.453	0.147	0.222
Shoubu B	8	0.327	0.018	0.132
	32	0.344	0.020	0.181
	128	0.348	0.027	0.132
	512	0.360	0.030	0.135

Table 4. Breakdown of calculation time for strong-scaling runs

System	# processes	interaction	comm.	others
TaihuLight	10000	2.0738	0.0410	0.4658
	20000	1.0499	0.0253	0.2426
	40000	0.5565	0.0298	0.1125
	80000	0.2991	0.0233	0.0652
	160000	0.1765	0.0322	0.0356
Gyoukou	1024	0.3323	0.1140	0.2808
	2048	0.1512	0.0668	0.1658
	4096	0.0854	0.0417	0.0923
	8192	0.0538	0.0357	0.0582

number of particles again, the actual value would go back to 300 since we go down the tree one level. Thus, the size of particle group to share the interaction list shows oscillation as the function of the total number of particles, and that is the reason for the oscillation of the calculation time.

Figures 5 shows the strong scaling result. The total number of particles is 10^{11} and 10^{10} , for TaihuLight and Gyoukou. We do not show the strong-scaling result for Shoubu System B, since it is rather small system and strong-scaling result is not so meaningful. We can see that speedup is almost linear.

Tables 3 and 4 show the breakdown of the calculation time per one timestep, again for both the weak and strong scaling runs. As expected, in the case of strong-scaling runs, the calculation time for communication does not decrease significantly, and eventually limits the performance. As already stated, our main interest is to use very large number of particles. Therefore, for actual scientific runs, the communication time would not become the limiting factor.

If we compare the calculation times on Gyoukou and Shoubu System B, we can see that the times for the interaction calculation are similar. but for both communications and "others", Shoubu System B is much faster. Again, this is not due to hardware difference but due to software difference.

The performance of run for 1.6×10^{12} particles on 160k processes (40000 nodes) of TaihuLight is 47.9 PF, or 39.7% of the theoretical peak performance of the Sunway TaihuLight system. On PEZY-SC2 based systems, we achieved 10.6PF for 8×10^9 particles on 8K SC2 chips, or efficiency of 23.3% of the theoretical peak performance.

On 512-chip Shoubu System B, we achieved the speed of 1.01 PF, or 35.5% §.

The overall efficiency we achieved on PEZY-SC2 based systems is a bit lower compared to that on TaihuLight. This difference is not due to any fundamental difference in the architecture but purely due to the limitation on the available time for program development and performance measurement. As we stated, the calculation in the list-construction step, such as the constructions of the tree and the interaction lists are currently done on Xeon D, and around 40% of the total time is consumed in this step at the time of measurement on Gyoukou. Most of these are now done on SC2 side, and that is why the performance of Shoubu System B is better than that of Gyoukou.

In terms of the number of particles integrated per second, we have achieved 5.5×10^{11} particles per second, which is more than 10 times faster than the results of previous works on K computer (Ishiyama et al. 2012) or ORNL Titan (Bédorf et al. 2014). In other words, on modern, "bigger" and low B/F supercomputers, we can achieve the performance efficiency almost same as that have been achieved on old, "smaller" and high B/F supercomputers, by using our algorithms.

6 Discussion and summary

6.1 Midpoint Tree Construction Method

As we have described in section 4.2, the persistent interaction list method gives the additional force error from the deformation of the tree cells due to shear motion. We found that the time integration error grows rather rapidly because the error due to shear deformation of tree cells grows linearly in time. As a result, the error, such as the error of the total energy, grows as square of time, while the same interaction list is used. Figure 6 shows the evolution of relative energy error. The definition of total energy is given by equation 3.

$$\begin{aligned}
 E &= E_{\text{kin}} + E_{\text{pot}} + E_{\text{disp}}, \\
 E_{\text{kin}} &= \sum_i \frac{m_i v_i^2}{2}, \\
 E_{\text{pot}} &= \sum_{i < j} \phi_{ij} - \sum_i \frac{G m_i m_p}{r_{ip}}, \\
 E_{\text{disp}} &= - \sum_i \int \left(\sum_j \mathbf{F}_{ij, \text{disp}} \right) \cdot \mathbf{v}_i dt, \\
 \phi_{ij} &= \begin{cases} \frac{-G m_i m_j}{r_{ij}} & (r_{ij} > r_{\text{coll}}) \\ \frac{m_i m_j}{2} \left(\frac{G r_{ij}^2}{r_{\text{coll}}^3} - \frac{3G}{r_{\text{coll}}} - \frac{\kappa (r_{ij} - r_{\text{coll}})^2}{m_i + m_j} \right) & (r_{ij} \leq r_{\text{coll}}) \end{cases}, \\
 \mathbf{F}_{ij, \text{disp}} &= \begin{cases} \mathbf{0} & (r_{ij} > r_{\text{coll}}) \\ - \frac{\eta m_i m_j \mathbf{r}_{ij} \cdot \mathbf{v}_{ij}}{(m_i + m_j) r_{ij}^2} \mathbf{r}_{ij} & (r_{ij} \leq r_{\text{coll}}) \end{cases},
 \end{aligned}
 \tag{5}$$

where E is the total energy, E_{kin} is the kinetic energy, E_{pot} is the potential energy for gravitational and restitution forces and E_{disp} is the energy dispated by particle-particle collision.

We can clearly see that the error with the persistent interaction list method (solid curve) grows faster than that without the method (dashed curve).

To reduce this integration error, we developed a new method




Figure 6. Energy error as a function of time. Solid, dashed and dotted curves indicate the result with the persistent interaction method (PI), without PI and with PI and the midpoint tree construction method (MT). Bottom panel is zoom-in of the top panel.

in which the interaction list is constructed at the midpoint of the period for which the same interaction list is used. For example, if we use the same interaction list for 64 steps, we construct the interaction list by using the predicted positions of particles at 32 steps after the current time. With this algorithm, the shear deformation of cells in first half of the reusing period and that in the last half are in opposite directions, and we can expect that the dominant error term would cancel out. We implemented this method to our code and confirmed that this algorithm actually cancels the dominant error. As a result, by using this method we do not have to make timestep or opening criterion small.

Unfortunately, we did not use this method for the performance measurements. However, the additional cost is very small because particles are almost on the circular orbits and the prediction is easy. Thus even if we use this method, the performance results would not change.

(8)

§If we use the number of interactions when the interaction list is not shared for multiple particles (i.e. the number of interactions is minimum) to measure the performance, the number of interactions is reduced by 25% and the performances are 36, 8 PF and 0.75 PF on TaihuLight, Gyoukou and Shoubu System B.

6.2 Performance Portability

We have reported the measured performance of two rather different HPC systems, Sunway TaihuLight and PEZY-SC2 based systems, for the same large-scale simulation of self-gravitating planetary rings. In both cases, we have achieved fairly high efficiency, more than 30% of the theoretical peak. The parallel algorithm used is essentially the same for the two systems. However, the actual codes are rather different, simply because of the difference in the architecture and the software development environment.

Sunway TaihuLight has a heterogeneous many-core architecture integrated in one chip. Thus, the CPU (MPE in their terms) and accelerators (CPE in their terms) share the same physical memory, but CPEs lack the data cache and need to rely on DMA controller to access the main memory efficiently.

On TaihuLight, one can use OpenACC compiler. However, in order to achieve high performance, one is practically forced to use the Athread[¶] call, which makes the 64 CPEs and their local memories visible to programmers.

On the other hand, PEZY-SC2 systems, at least at present, have a rather classical accelerator-based architecture, in which CPU (a Xeon-D processor) and accelerators (PEZY-SC2 processors) are connected through PCI Express interface. This means that they have separate physical memories. Within one chip, however, processing elements of PEZY-SC2 processor have three levels of data caches. Currently PZCL, a dialect of OpenCL, is supported on PEZY-SC2 based systems.

Because of these differences (shared and separate memory, DMA and cache, thread-based and OpenCL-like), the actual programs for two machines have become quite different, even though the algorithms used are the same and the problem to be solved is the same.

Both codes, however, are based on our framework, FDPS (Iwasawa et al. 2016; Namekata et al. 2018) and follow its structure. The basic idea of FDPS is to separate the implementation of parallel algorithms and description of the physical problem. FDPS provides the former and the application programmers provides the latter, in the form of the data type definition of particles and functional form of particle-particle interaction. Users of FDPS can write their programs by specifying the data structure of particles they use, and calling necessary FDPS functions for domain decomposition, particle migration between processes, and interaction calculation. Currently, users should provide optimized function for particle-particle interaction calculation.

Many of the parallel algorithms we newly implemented are not specific to planetary rings but can be applied to any other particle-based simulations. Using FDPS, users can write their programs in their favorite language (currently C++, Fortran and C are supported) (Namekata et al. 2018), and let FDPS do complex parallelization.

Thus, it seems that one way to achieve performance and program portability on new machines with rather exotic architecture such as the machines evaluated in this paper is to develop the framework with a common API and internal implementations specialized and optimized to specific architectures. This is fairly straightforward in the case of TaihuLight, in which the CPE and MPEs share the

single physical memory, since the data structure that FDPS handles can still in the shared main memory. The basic data structure of FDPS is just an array of particles, and both the user-developed application program and the FDPS side can access that particle array in usual way.

On the other hand, how the separate memory spaces of PEZY-SC2 should be handled within FDPS requires a bit more consideration. One possibility would be to add the interface in which the user-side programs, for example the function to perform time integration, is passed to FDPS, instead of directly called within the user-side program. Here, the function to be passed applies to single particle (or some small array of particles), and applying it to all particles in the system will be the responsibility of FDPS. This approach will probably make the software development and performance improvement easier on other machines, since parallelization in both MPI and OpenMP level can be taken care of within FDPS.

This problem of portability is of course not limited to FDPS. It occurs in practically any application in any field of computational science. We clearly need a new and systematic approach to solve this problem, and we think the use of frameworks such as FDPS may be an efficient and practical way.

Our view of frameworks is that they should allow users to express their problems in simple and (ideally) machine-independent way. In the case of particle-based simulations, we have designed FDPS to meet this goal, and it actually works pretty well on large HPC systems, both with and without GPGPUs. Our experience on TaihuLight and PEZY-SC2 indicates that it is also possible to extend FDPS to cover these systems. We believe similar approaches will be used in other fields.

6.3 Summary

In this paper, we described the implementation and performance of a highly efficient simulation code for self-gravitating planetary rings on Sunway TaihuLight and PEZY-SC2 based systems.

The measured performance is 47.9 PF, or 39.7% of the theoretical peak, for simulation of 1.6×10^{12} particles on 40,000 nodes of TaihuLight, 10.6PF, or 23.3% of the theoretical peak, for simulation of 8×10^{10} particles on 8192 nodes of Gyoukou, and 1.01PF, or 35.5% of the theoretical peak for 5×10^9 particles on 512 nodes of Shoubu System B. As noted earlier, Gyoukou and Shoubu System B use the same PEZY-SC2 processor. The difference in the efficiency is purely due to the fact that Gyoukou was turned off on March 31, 2018. The software at that time was still under development.

Compared to previous achievements on K computer or ORNL Titan, the achieved efficiency is similar or higher, and the speed in terms of the number of particles integrated per second is higher, for both TaihuLight and PEZY-SC2 based systems. As we stated earlier, this level of performance would not be achieved without the new algorithms described in this paper.

[¶]Atherad is a thread library to use CPEs in parallel.

Compared to other multi-core processors for modern HPC systems such as Fujitsu SPARC64 VIIIfx and IXfx or Intel Xeon Phi processors, both SW26010 processor of TaihuLight and PEZY-SC2 processor of PEZY-SC2 based systems have several unique features which allow very high peak performance but at the same time make it much harder to achieve high efficiency on real applications. These are:

- Heterogeneous architecture with rather extreme performance ratio of 1:64 in the case of SW26010 and even larger in the case of SC2.
- The lack of cache hierarchy (SW26010) or cache coherency (SC2).
- Very limited main memory bandwidth, with B/F values around 0.02–0.04. This is about 1/10 of the numbers of Fujitsu or Intel HPC processors.

On the other hand, SW26010 comes with very well-thought features which allows the programmers to optimize the performance of code on CPE. These features include:

- Low-latency DMA controller which can be initiated by any CPE.
- Low-latency, high-bandwidth communication between CPEs.

These two features allow very efficient use of the main memory bandwidth. The two-dimensional structure of the network within CG seem to be optimized for highly efficient implementation of matrix-matrix multiplications, but it is actually quite useful for other real applications, whenever fast inter-core communication is necessary.

It is certainly true that the need to use DMAs for data transfer between CPE and main memory complicates the use of CPE. However, it is also true that it makes quite optimized access to main memory possible, since the application programmer can (or have to) control all main memory accesses. In the case of our code, in several places we have “vectorizable” loops, which perform the same operation on all particles in the system. The number of operations per particle is relatively small, of the order of ten, and the data size of one particle is 32 bytes. In the case of manycore architecture with hierarchical cache memory, to achieve high efficiency on simple vector operations like

```
a[i] = b[i]+c[i]
```

is actually quite complicated. In modern processors, load address would be predicted and hardware prefetch is generated. The hardware prefetch would probably work for a very simple loop like the above example, but would fail if many vectors are loaded. Then the programmer need to experiment with software prefetch, to find the way to get the best performance.

In the case of SW26010, currently it is rather tedious and error-prone to write the equivalent operation using the combination of Athread and DMA, and sometimes inner kernel in assembly language, but once we do so, we can get a performance close to the theoretical limit relatively easily.

The existence of low-latency (less than 10 clock cycles) communication path between CPEs is quite important for using CPEs for fine-grain parallelism such as loop-level parallelization. Such low-latency communication is difficult to implement on shared memory processors with hierarchical cache.

The SC2 processor supports the cache flush and synchronization at each level of the cache hierarchy, making the relatively low-latency communications between processors possible. However, it is clearly desirable to have more direct control of interprocessor communication.

One common problem of SW26010 or SC2 is that writing high-performance kernel for them means writing the inner kernel in the assembly language. This is purely the software limitation, and probably not so difficult to fix. In this aspect, SC2 is somewhat easier to deal with, since it supports 8-way multithreaded execution, which effectively hides the latencies of L1 cache and arithmetic unit from compiler.

In conclusion, we have implemented parallel particle simulation code on Sunway SW26010 and PEZY-SC2 processors, and found that it is not impossible to achieve high performance on their rather extreme architectures with carefully designed algorithms. Even though the B/F number are less than 0.1 and the network bandwidth is similarly low, the efficiency we have achieved is comparable to that on K computer, with 15 times more memory and network bandwidth. We feel that architecture evolution in this direction will help the HPC community to continue improving the performance. We also believe that high-level software framework such as our FDPS will help many researchers to run their own applications efficiently on new architectures.

Acknowledgment

We are deeply indebted to anonymous referees for valuable comments. Numerical computations were in part carried out on K computer at RIKEN Center for Computational Science through the HPCI System Research project (Project ID:ra000008), on HOKUSAI at RIKEN Advanced Center for Computing and Communication, and on Cray XC50 at Center for Computational Astrophysics, National Astronomical Observatory of Japan. This work was supported by The Large Scale Computational Sciences with Heterogeneous Many-Core Computers in grant-in-aid for High Performance Computing with General Purpose Computers in MEXT (Ministry of Education, Culture, Sports, Science and Technology-Japan), by JSPS KAKENHI Grant Number JP18K11334 and JP18H04596 and also by JAMSTEC, RIKEN, and PEZY Computing, K.K./ExaScaler Inc. In this research computational resources of the PEZY-SC2 based systems supercomputer, developed under the Large-scale energy-efficient supercomputer with inductive coupling DRAM interface project of NexTEP program of Japan Science and Technology Agency, has been used.

References

- Barnes J and Hut P (1986) A hierarchical o(nlogn) force calculation algorithm. *Nature* 324: 446–449.
- Barnes JE (1990) A modified tree code: don’t laugh; it runs. *Journal of Computational Physics* 87: 161–170.
- Bédorf J, Gaburov E, Fujii MS, Nitadori K, Ishiyama T and Zwart SP (2014) 24.77 pflops on a gravitational tree-code to simulate the milky way galaxy with 18600 gpus. In: *SC14: International Conference for High Performance Computing, Networking, Storage and Analysis*. pp. 54–65. DOI:10.1109/SC.2014.10.

- Blelloch GE, Leiserson CE, Maggs BM, Plaxton CG, Smith SJ and Zagha M (1991) A comparison of sorting algorithms for the Connection Machine CM-2. In: *Proceedings Symposium on Parallel Algorithms and Architectures*. Hilton Head, SC, pp. 3–16.
- Cai L, Wang YC, Tang W, Wang B, Ethier S, Liu Z and Lin J (2018) Openacc vs the native programming on sunway taihulight: a case study with gtc-p. In: *Proceedings of the IEEE International Conference on Cluster Computing*. IEEE, pp. 88–97.
- Fu H, Liao J, Yang J, Wang L, Song Z, Huang X, Yang C, Xue W, Liu F, Qiao F, Zhao W, Yin X, Hou C, Zhang C, Ge W, Zhang J, Wang Y, Zhou C and Yang G (2016) The sunway taihulight supercomputer: system and applications. *Science China Information Sciences* 59(7): 072001. DOI:10.1007/s11432-016-5588-7. URL <http://dx.doi.org/10.1007/s11432-016-5588-7>.
- Hamada T, Narumi T, Yokota R, Yasuoka K, Nitadori K and Taiji M (2009) 42 tflops hierarchical n-body simulations on gpus with applications in both astrophysics and turbulence. In: *Proceedings of the Conference on High Performance Computing Networking, Storage and Analysis*, SC '09. New York, NY, USA: ACM. ISBN 978-1-60558-744-8, pp. 62:1–62:12. DOI:10.1145/1654059.1654123. URL <http://doi.acm.org/10.1145/1654059.1654123>.
- Ishiyama T (2014) Hierarchical Formation of Dark Matter Halos and the Free Streaming Scale. *The Astrophysical Journal* 788: 27. DOI:10.1088/0004-637X/788/1/27.
- Ishiyama T, Nitadori K and Makino J (2012) 4.45 pflops astrophysical n-body simulation on k computer: The gravitational trillion-body problem. In: *Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis*, SC '12. Los Alamitos, CA, USA: IEEE Computer Society Press. ISBN 978-1-4673-0804-5, pp. 5:1–5:10. URL <http://dl.acm.org/citation.cfm?id=2388996.2389003>.
- Iwasawa M, Tanikawa A, Hosono N, Nitadori K, Muranushi T and Makino J (2016) Implementation and performance of FDPS: a framework for developing parallel particle simulation codes. *Publications of the Astronomical Society of Japan* 68: 54. DOI: 10.1093/pasj/psw053.
- Makino J (1991) Treecode with a special-purpose processor. *Publications of the Astronomical Society of Japan* 43: 621–638.
- Makino J (2004) A Fast Parallel Treecode with GRAPE. *Publications of the Astronomical Society of Japan* 56: 521–531.
- McMillan SLW and Aarseth SJ (1993) An $O(n \log n)$ integration scheme for collisional stellar systems. *The Astrophysical Journal* 414: 200–212.
- Michikoshi S and Kokubo E (2017) Simulating the Smallest Ring World of Chariklo. *The Astrophysical Journal Letters* 837: L13. DOI:10.3847/2041-8213/aa6256.
- Namekata D, Iwasawa M, Nitadori K, Tanikawa A, Muranushi T, Wang L, Hosono N, Nomura K and Makino J (2018) Fortran interface layer of the framework for developing particle simulator FDPS. *Publications of the Astronomical Society of Japan* 70: 70. DOI:10.1093/pasj/psy062.
- Portegies Zwart S and Bédorf J (2014) Computational Gravitational Dynamics with Modern Numerical Accelerators. *ArXiv e-prints*.
- Rein H and Latter HN (2013) Large-scale N-body simulations of the viscous overstability in Saturn's rings. *Monthly Notices of the Royal Astronomical Society* 431: 145–158. DOI:10.1093/mnras/stt152.
- Rein H and Liu SF (2012) REBOUND: an open-source multi-purpose N-body code for collisional dynamics. *Astronomy & Astrophysics* 537: A128. DOI:10.1051/0004-6361/201118085.
- Salmon J, Quinn PJ and Warren M (1990) *Using Parallel Computers for Very Large N-Body Simulations: Shell Formation Using 180 K Particles*. Berlin, Heidelberg: Springer Berlin Heidelberg. ISBN 978-3-642-75273-5, pp. 216–218. DOI:10.1007/978-3-642-75273-5_51. URL http://dx.doi.org/10.1007/978-3-642-75273-5_51.
- Warren MS and Salmon JK (1992) Astrophysical N-body simulations using hierarchical tree data structures. In: *Supercomputing '92*. Los Alamitos: IEEE Comp. Soc., pp. 570–576.
- Wisdom J and Tremaine S (1988) Local simulations of planetary rings. *The Astronomical Journal* 95: 925–940. DOI:10.1086/114690.
- Zebker HA, Marouf EA and Tyler GL (1985) Saturn's rings: Particle size distributions for thin layer models. *Icarus* 64(3): 531 – 548. DOI:[http://dx.doi.org/10.1016/0019-1035\(85\)90074-0](http://dx.doi.org/10.1016/0019-1035(85)90074-0). URL <http://www.sciencedirect.com/science/article/pii/0019103585900740>.

Author Biographies

Masaki Iwasawa received PhD from the University of Tokyo. After he received PhD, he worked at the National Astronomical Observatory of Japan, Leiden University, RIKEN and Kobe University. Since April 2020, he works at National Institute of Technology. His research interests are computational astronomy and high-performance computing.

Daisuke Namekata received PhD from Hokkaido University. After he received PhD, he worked at the National Astronomical Observatory of Japan and the University of Tsukuba. Since March 2016, he works at RIKEN Center for Computational Science in Particle Simulator Research team. His research interest are physics of interstellar medium, radiation hydrodynamic simulation, and high-performance computing.

Ryo Sakamoto is software engineer work at PEZY Computing K.K. His interests are high-performance computing and embedded systems technology.

Takashi Nakamura received his Associate degree from Osaka Prefecture University. He is presently a software engineer at Preferred Networks, Inc. He develops software for high-performance computing and framework for hardware accelerators.

Yasuyuki Kimura is software engineer at PEZY Computing K.K. His interests are computer science, electrical engineering and embedded systems technology.

Keigo Nitadori received PhD from University of Tokyo in 2009. After he received PhD, he worked at RIKEN and Tsukuba University. Since 2019, he is a technical scientist of RIKEN R-CCS.

Long Wang received astrophysical PhD from the Peking University at 2017. After that, he works as a postdoctoral researcher at RIKEN-CCS, Kobe, Japan. In 2018, he got the Alexander von Humboldt postdoctoral fellowship and worked at the University of Bonn, Germany. Now he is a JSPS postdoctoral fellowship

working at the University of Tokyo. His research interest is high-performance computing and dynamical evolution of astronomical N-body systems.

Miyuki Tsubouchi joined Particle Simulator Research team of RIKEN Advanced Institute of Computational Science (AICS, now R-CCS) in January 2014. Since then, she works as a technical staff, and since April 2014 she also serves as a member of Co-design team of Flagship 2020 project within AICS/R-CCS.

Jun Makino received PhD from the University of Tokyo. After he received PhD, he worked at University of Tokyo, the National Astronomical Observatory of Japan, and Tokyo Institute of Technology. Since April 2014, he is a subleader of the exascale computing project and the team leader of the Co-design team, AICS, RIKEN, and since March 2016 he works also at Kobe University. His research interests are stellar dynamics, large-scale scientific simulation and high-performance computing.

Zhao Liu is a PhD candidate in computer science at Tsinghua University and a researcher at National Supercomputing Center in Wuxi. His research interests include high performance computing and computer architecture.

Haohuan Fu is a professor in the Ministry of Education Key Laboratory for Earth System Modeling, and Department of Earth System Science in Tsinghua University, where he leads the research group of High Performance Geo-Computing (HPGC). He is also the deputy director of the National Supercomputing Center in Wuxi, leading the research and development division. Fu has a PhD in computing from Imperial College London. His research work focuses on providing both the most efficient simulation platforms and the most intelligent data management and analysis platforms for geoscience applications, leading to two consecutive winning of the ACM Gordon Bell Prizes (nonhydrostatic atmospheric dynamic solver in 2016, and nonlinear earthquake simulation in 2017).

Guangwen Yang is a professor in the Department of Computer Science and Technology, and Beijing National Research Center for Information Science and Technology, at Tsinghua University, China. He is also the director of the National Supercomputing Center in Wuxi, China. His research interests include parallel algorithms, cloud computing, machine learning, and the earth system model. Yang has a PhD in computer science from Harbin Institute of Technology. He has been awarded the ACM Gordon Bell Prize (2016, 2017). He is a member of IEEE.