

# Portable performance through frameworks: Experience on extreme-scale heterogeneous many-core architectures

Masaki Iwasawa\*, Daisuke Namekata\*, Ryo Sakamoto<sup>†</sup>, Takashi Nakamura<sup>†</sup>, Yasuyuki Kimura<sup>‡</sup>, Keigo Nitadori\*, Long Wang<sup>§</sup>, Miyuki Tsubouchi\*, Jun Makino<sup>¶</sup>, Zhao Liu<sup>||\*\*</sup>, Haohuan Fu<sup>††</sup>, Guangwen Yang<sup>‡‡</sup>

\*RIKEN Center for Computational Science

7-1-26 Minatojima-minami-machi, Chuo-ku, Kobe, Hyogo, 650-0047, Japan

Email: masaki.iwasawa@riken.jp

<sup>†</sup> PEZY Computing K. K., 1-11 Ogawa-machi, Chiyoda-ku, Tokyo, 101-0052, Japan

Email: yasuyuki@pezy.co.jp

<sup>‡</sup> ExaScaler Inc., 2-1 Ogawa-machi, Chiyoda-ku, 101-0052 Tokyo, Japan

Email: yasuyuki@exascalr.co.jp

<sup>§</sup> Universität Bonn, Argelander-Institut für Astronomie, Auf dem Hügel 71, 53121 Bonn

<sup>¶</sup> Kobe University, 1-1, Rokkodai-cho, Nada-ku, Kobe, 657-8501, Japan

Email: jmakino@riken.jp

<sup>||</sup> Department of Computer Science and Technology, Tsinghua University, Beijing, China.

<sup>\*\*</sup> National Supercomputing Center in Wuxi, Wuxi, China

Email: liuzhao@mail.nscwx.cn

<sup>††</sup> Department of Earth System Science, Tsinghua University, Beijing, China

Email: haohuan@tsinghua.edu.cn

<sup>‡‡</sup> Department of Computer Science and Technology, Tsinghua University, Beijing, China

Email: ygw@tsinghua.edu.cn

**Abstract**—In this paper, we report implementations and measured performance of our extreme-scale global simulation code for planetary rings on Sunway TaihuLight and two PEZY-SC2 systems: Shoubu System B and Gyoukou. The numerical algorithm is the parallel Barnes-Hut tree algorithm, which has been used in many large-scale astrophysical particle-based simulations. Our implementation is based on our FDPS framework. However, the extremely large numbers of cores of the systems used (10M on TaihuLight and 16M on Gyoukou) and their relatively poor memory and network bandwidth pose new challenges. We describe the new algorithms introduced to achieve high efficiency on machines with low memory bandwidth. The measured performance is 47.9, 10.6 PF, and 1.01PF on TaihuLight, Gyoukou and Shoubu System B (efficiency 40%, 23.5% and 35.5%).

## I. INTRODUCTION

Saturn’s ring was first observed by Galileo Galilei in 1610. For more than three centuries, it had been the only known ring system within our solar system. In 1977, rings of Uranus were found through occultation observations from an aircraft, and then in 1979 rings of Jupiter by Voyager 1 and in 1989 those of Neptune by Voyager 2. Very recently, it turned out that some of minor planets also have rings. The first distinctive example is 10199 Chariklo, whose orbit is between those of Saturn and Uranus (and thus one of Centaurs). There are probably more Centaurs with rings.

Thus, quite recently, a wide variety of ring systems have been found. How these rings were formed and have evolved is an important question in planetary science, and large-scale, global simulation, if possible, would help greatly to advance our understanding.

Planetary rings are usually at the radii around the Roche limit. Thus, mutual gravity between particles does not easily lead to the formation of new satellites, but is important enough to form spiral waves (“wakes”) in very small scales, which increase the effective viscosity and should enhance the radial transport of the angular momentum. On the other hand, the actual ring system seems to consist of very large number of narrow rings, separated with distinct gaps. It is believed that these gaps are maintained by high-order resonances with small embedded satellites (so-called moonlets), but whether or not such gaps can be formed and maintained by resonances has not been fully understood.

Up to now, most of simulations of ring structures have been local ones, in which a small patch was cut out from the ring and simulated under the assumption of the local Hill approximation and periodic boundary condition [?]. Rein and Latter [?] performed “Large-scale” simulation of viscous overstability in Saturn’s rings, using up to 204,178 particles and up to 10,000 orbits using this local approach. Because very long simulations are necessary, the number of particles

has been small. They used REBOUND [?], an MPI-parallel  $N$ -body simulation code.

Michikoshi and Kokubo [?] performed global simulations of rings with the largest number of particles reported so far. They used 300M particles to model two narrow rings of Chariklo. They have developed their parallel code using the framework we developed, FDPS [?].

Almost all previous studies of planetary rings adopted so-called “local” approximation, in which only a small patch of a ring is simulated assuming periodic boundary condition in both radial and azimuthal directions.

Michikoshi and Kokubo [?] performed global simulations of rings with 300M particles, using FDPS [?]. They so far followed the system only for 10 orbital periods.

The total calculation cost is roughly proportional to number of particles multiplied by the number of orbital periods followed, since the calculation cost per timestep is  $O(N \log N)$  when Barnes and Hut tree algorithm is used and the number of timestep required for ring simulations is essentially independent of the number of particles. Thus, we can conclude that the size of state-of-the-art simulations of planetary rings is around  $10^9$  particle-orbits, or around  $10^{12}$  particle-steps.

We should note that even though the simulations so far done in this field is relatively small, that does not mean there is no need or possibilities for larger scale simulations. If we want to model the global structures of rings, we cannot rely on local treatment. For example, the effect of resonances with small satellites can only be studied using global simulations. On the other hand, the number of particles one need for global simulations, even for a very narrow radial range, is very large. For example, consider the A ring of Saturn with the radius of around  $1.3 \times 10^5$  km. The typical radius of ring particles is 6 m [?], and the optical depth of the ring is around unity. Thus, we need  $10^4$  particles per square km or around  $10^{12}$  particles for the radial range of 100 km. With this radial range, we can model many of fine features observed by Cassini directly.

If we could use particles with larger size, we could reduce the number of particles required significantly. However, that would change the viscous diffusion timescale of the ring, and thus what would be observed. It is necessary to perform simulations with particles of real physical radius, which would require at least  $10^{16}$  and ideally  $10^{19}$  particle steps.

In other fields of astrophysics, very large simulations have been performed. For example, Ishiyama [?] used  $4096^3$  particles to follow the formation and growth of dark matter halos of smallest scales. This simulation corresponds to  $10^{16}$  particle steps. Part of this calculation was performed on K computer. The performance of K computer is  $4.0 \times 10^{10}$  particle steps per second on the entire K computer, or 60,000 particle step per second per core for a processor core with the theoretical peak performance of 16 Gflops [?]. The efficiency they achieved is 55% of the theoretical peak.

The algorithms used in large-scale  $N$ -body simulations are rather similar, and that means they are well studied and close to optimal. All of them use domain decomposition and Barnes and Hut tree algorithm. For domain decomposition, several

variations have been used, such as Orthogonal Recursive Bisection [?], Hashed Oct Tree [?], Multisection [?].

Efficient implementations on large-scale GPGPU clusters exist [?], [?], [?]. Bédorf *et al.* performed the simulation of Milky Way Galaxy using  $2.42 \times 10^{11}$  particles. The achieved performance is 24.77 PF on ORNL Titan, and one timestep took 5.5 seconds. Thus they have achieved the performance of  $4.4 \times 10^{10}$  particle steps per seconds. The theoretical peak performance of Titan is 73.2 PF in single precision. Thus, the achieved efficiency is 33.8%.

So far, there is no report on the implementation and performance of parallel tree algorithm on processors with a heterogeneous architecture such as Sunway TaihuLight [?] or systems based on PEZY-SC2 processors. In this paper, we present the result of our effort to implement the parallel tree algorithm on them.

Simulation programs for two architectures are both based on our FDPS framework for large-scale parallel particle-based simulations, and we implemented essentially the same algorithms to improve the parallel efficiency for both systems. Currently, the actual codes are different, since Sunway TaihuLight and PEZY-SC2 systems have quite different parallel programming environments. However, we are working on new version of FDPS, which allows the user program runs on these systems and also on usual CPU- or GPGPU-based systems, with minimal changes of the application program. In the rest of this paper, we first give a short description of the architecture of the two systems in section II. In section III, we describe in detail the new algorithms we developed to achieve high performance on extreme-scale heterogeneous many-core architectures. In section IV we describe how the performance was measured and achieved performance. In section V, we discuss what would be necessary to achieve performance portability over wide range of hardware architectures and parallel programming environments. Section VI is for summary.

## II. SUNWAY TAIHULIGHT AND PEZY-SC2 SYSTEMS

In this section, we briefly describe the features of the Sunway TaihuLight system and two systems with PEZY-SC2 processors: Gyoukou and Shoubu System B. For more details of TaihuLight system see [?]. TaihuLight consists of 40960 Sunway 26010 processors, and Gyoukou and Shoubu System B 13312 and 512 PEZY-SC2 processors, respectively. Unfortunately, Gyoukou was turned off by March 31, 2018, and thus our performance measurement on Gyoukou system was based on a preliminary version of the simulation code, and thus the efficiency is lower than that measured on Shoubu System B.

One SW26010 processor consists of four CGs (core groups), each with one MPE (management processing element) and 64 CPEs (computing processing elements). Both MPE and CPE are 64-bit RISC cores. MPE has L1 cache memories for both instructions and data, and also L2 data cache. On the other hand, each CPE has L1 instruction cache and 64KB of local data memory. CPEs can communicate with the main memory

through DMA. Each CPE can initiate multiple asynchronous DMA operations.

Each core group is connected to 8GB DDR3 DRAM memory with the theoretical peak transfer rate of 34GB/s. The processor runs at the clock frequency of 1.45GHz, and each core (both MPE and CPE) can perform four double precision FMA operations. Thus, the theoretical peak performance of one processor is 3016 Gflops and that of one CG is 754 Gflops. Thus, even when we use the nominal number for DRAM bandwidth, the B/F ratio is only 0.045. This is less than 1/10 of the number for K computer.

Compared to that of K computer, the network is also weak, with the total bandwidth of around 10GB/s per node. This is about the same as the performance of a single link of 6D torus network of K computer. Since the SW processor is around 25 times faster than the SPARC64 processor of K computer, the relative network bandwidth is different by more than two orders of magnitudes.

One PEZY-SC2 processor chip consists of 2048 processors (64 of them are disabled and the available number of processors is 1984). Each processor can perform 1, 2 and 4 multiply-and-add operation for FP64, FP32, and FP16 data. For FP32 and FP16, 2- and 4-way SIMD operations are performed. With the clock speed of 700MHz, the theoretical peak speed is 2.8, 5.6 and 11.1TF, for FP64, FP32 and FP16, respectively. At present, each SC2 processor chip have 4 channels of DDR4 memory, for the peak throughput of 76.8GB/s. Thus B/F is 0.027.

They have three levels of shared cache, but without coherency. Instead, they have explicit cache flush instructions to each levels. Two processors share L1D, and 16 processors L2D, and all processors LLC. Each processor runs either four or eight threads simultaneously. Thus, it is relatively easy to hide the latency of the arithmetic units and L1D.

In the original design, each SC2 processor chip had six MIPS64 cores, which were supposed to run the operating system and main body of the application programs. Unfortunately, currently they are disabled, and operating system and application programs run on the frontend Xeon D-1571 processor. Each Xeon D hosts eight SC2 processors. Thus, the performance ratio between Xeon D and SC2 is close to 100. Moreover, these eight SC2 are connected to Xeon D through single PCIe Gen3 16-lane channel. Thus, the peak data transfer speed between one SC2 and Xeon D is 2 GB/s, for the peak speed of 2.8TF.

In summary, TaihuLight and two systems based on PEZY-SC2 processors share the following characteristics:

- 1) Very large performance ratio between “general-purpose” and “computing” cores, close to 1:100.
- 2) Very small memory B/F numbers, around 0.03.
- 3) Even smaller network B/F numbers, 0.006 or 0.001.
- 4) Very large number of MPI processes, 160k or 10k.
- 5) Very large number of “computing” cores per MPI process, 64 or 1984.

Just one of these characteristics makes it very difficult to achieve reasonable performance for particle-based simulations

using previously known parallelization algorithms. In the next section, we describe the new algorithms we implemented to achieve good performance on these systems.

### III. NEW ALGORITHMS FOR EXTREME-SCALE SIMULATIONS

#### A. Overview of new algorithms

In this section, we describe the new algorithms we made in order to utilize TaihuLight and PEZY-SC2 based systems for the simulations of self-gravitating planetary rings. The following is the list of new algorithms.

- 1) The reuse of the interaction list over multiple timesteps.
- 2) Elimination of the global all-to-all communication.
- 3) “Semi-dynamic” load balance between computing cores
- 4) Optimizations specific to the ring geometry.

For TaihuLight and PEZY-SC2 based systems, we have modified our FDPS framework in architecture-specific way so that we implement the algorithms and run the code under the limited available time and software environment. However, many of these algorithms are ported back to the original FDPS so that anybody who uses FDPS can take advantage of these new algorithms.

In the rest of this section, we briefly describe these new innovations.

#### B. The reuse of the interaction list

The following gives the usual steps for highly parallel code for self-gravitating particle system:

- 1) Perform domain decomposition.
- 2) Exchange particles so that particles belong to appropriate domains.
- 3) Perform interaction calculation using fast algorithm such as Barnes-Hut tree.
- 4) Integrate the orbits of particles.
- 5) Go back to step 1.

In the case of approaches with local essential tree, step (3) consists of the following substeps:

- (3a) Construct the “local” tree structure from particles in the domain.
- (3b) Collect the information necessary for the calculation of interaction from other processes (so called local essential tree).
- (3c) Construct the “global” tree from the collected information.
- (3d) For small groups of particles, traverse the tree and calculate the interaction. Repeat this for all groups.

In the original algorithm [?], the traversal of the tree is done for each particle, and force calculation is done during the traversal. However, on almost all modern implementation, following the idea of Barnes [?], tree traversal are done for groups of neighboring particles, which are constructed using the tree structure itself. During the traversal for a group, the list of particles and tree nodes which exert the force on this group of particles is constructed, and actual force calculation

is done through the double loop over particles in the group and those in the list. This structure makes it possible to use vector pipelines, scalar SIMD units, and even special-purpose computers [?] with high efficiency. For GPGPUs, the extension of this algorithm, in which multiple lists are constructed and then sent to GPGPU, is used [?].

This approach does not work well on TaihuLight or PEZY-SC2 based systems, because of the low performance of general-purpose core and limited memory bandwidth. The performance we can achieve with either approach for ring simulation on these machines is less than 1%. Thus, it is necessary to reduce the cost of tree construction and tree traversal, and we achieved this by using the same interaction lists over multiple timesteps. We call this method the persistent interaction list method.

The idea behind this method is essentially the same as that for the neighbor-list method used in many simulation codes for particles with short-range interactions.

By using this persistent interaction list, we can reduce the calculation cost of the part other than the interaction calculation drastically. While we are using the same interaction lists, we skip the domain decomposition, exchange of particles, construction of the local tree. We still need to update the physical quantities of the nodes of the tree, since particles move at each timestep. We first update the information of the nodes of the local tree. Then, using the list of nodes for the local essential tree, the communication between the nodes is performed. Finally, the physical quantities of the global tree are updated, and the force calculation is performed using this updated global tree and the persistent interaction list.

The most time-consuming part of the tree construction is the sorting. In the case of TaihuLight, we implemented the parallel sample sort [?] on CPEs. In the case of Gyoukou, the sorting was performed on Xeon D host processor. In the case of Shoubu System B, it was performed on the side of PEZY-SC2 processors. Also, some other operations are moved from Xeon-D to PEZY-SC2. Thus, the overall performance is significantly better for Shoubu System B. As stated earlier, Gyoukou was turned off on March 31, 2018, and we could not measure the performance of our improved code on systems with more than 512 PEZY-SC2 processors.

We have ported all operations in timesteps in which the interaction list is used (list-reusing steps), except for MPI functions for communication, to CPEs (TaihuLight) or SC2 processors (PEZY-SC2 based systems). For the timestep in which the interaction list is constructed (list-constructing step), some of operations are still done on Xeon D in the case of PEZY-SC2 based systems.

### C. Tree and Domain structures on Cylindrical Coordinate

We want to model a relatively narrow ring, and this means the usual domain decomposition in Cartesian coordinates can cause serious problems. Figure 1 illustrate the problem. We can see the domains near the  $y$  axis are very elongated. This irregular shape of domains results in the increase of communi-




Fig. 1. Schematic figure of domain decomposition by the multisection method in  $x$ - $y$  coordinate. Domains are divided by  $16 \times 8$ .




Fig. 2. Schematic figure of domain decomposition by the multisection method in cylindrical coordinate. Domains are divided by  $4 \times 32$ .

cation between processes, and thus serious degradation in the efficiency.

We can avoid this problem, if we apply the domain decomposition in the cylindrical coordinates (figure 2). Note that we can also use the cylindrical coordinates for the construction of the tree. Since the ring is narrow, the local distance  $s$  in the Cartesian coordinates  $(x, y, z)$  can be approximated by that in cylindrical coordinates  $(r, \phi, z)$ .

$$ds^2 = dx^2 + dy^2 + dz^2 \sim d\phi^2 + dr^2 + dz^2, \quad (1)$$

when  $r \sim 1$ . Thus, we can use the cylindrical coordinate for domain decomposition and tree construction and even for the tree traversal, without any modification of the algorithm or program itself. The actual interaction calculation is faster in Cartesian coordinates and thus Cartesian coordinates is used.

### D. Coordinate rotation

The simulation of ring with very large number of processes poses new challenges. As we increase the number of processes, the size of the domains becomes smaller. On the other hand,

the timestep does not become much smaller even when we increase the total number of particles, since the random velocities of ring particles become smaller when we increase the number of particles. Thus, the distance that particles move can be comparable or even larger than the domain size, resulting in the increase in the amount of communication.

We can “solve” this problem by the rotation of the coordinates and domain structure, so that particles do not move much. If we rotate the coordinates at the speed of Kepler rotation at the center of the ring, particles at the center of the ring do not move much. Particles at other radial positions still move, but the speed becomes much smaller than that of the Kepler rotation. Thus, communication due to Kepler rotation can be almost eliminated.

#### E. Elimination of all-to-all communication

In FDPS, the exchange of LET (local essential tree) data is done though a single call to the `MPI_Alltoallv` function. This implementation works fine even for full-node runs on K computer, but becomes problematic on systems with relatively weak network like TaihuLight and PEZY-SC2 based systems. We can eliminate this all-to-all communication, by constructing the “tree of domains” locally and let only higher-level information be sent to distant processes.

In the current implementation specialized to narrow rings, we implemented a very simple two-level tree, in which the second-level tree nodes have all processes in the radial direction. For example, if we have a process grid of (1000, 10), where 1000 in angular and 10 in radial direction, 10 domains in the radial direction are combined to one tree node, resulting in 1000 second-level nodes. Only these 1000 nodes exchange their center-of-mass information. All LET information other than these center-of-mass data of second-level nodes are sent either to other second-level nodes (and then broadcast to lower-level nodes) or sent directly to lower-level nodes.

In this implementation, there is still one global communication in the angular direction, but we can use `MPI_Allgather` since only the top-level data are sent. Thus the reduction in the communication was quite significant.

#### F. Load Balance among computing cores

In our current implementation, interaction lists are created at the list-construction step, and are reused for several steps. The total number of lists in one MPI process is around  $10^5$ , and we need to use 64 or 1984 computing cores efficiently for them. If we just assign a fixed number of lists to cores, random variation of the list length can result in large load imbalance. Therefore, some load balance strategy is necessary. We applied the following simple algorithm.

- 1) Sort the interaction lists by their length.
- 2) Assign the longest 64 lists on 64 CPEs (in case of TaihuLight).
- 3) For each remaining list, assign it to the CPE with the shortest total calculation cost.

Since the calculation time of cores is quite predictable, this algorithm works very well.

In the case of PEZY-SC2 based systems, we further improved the load balance by using multiple cores which share the cache for one interaction list.

#### G. Interaction Kernel

In the case of TaihuLight, we found the compiler-generated code for the interaction kernel, even when SIMD operations are used, does not give very good performance. We rewrite the interaction kernel fully in the assembly language, with hand-unroll and careful manual scheduling. As a result, we achieved more than 50% of the theoretical peak performance for the kernel.

We have applied similar optimization also on PEZY-SC2 based systems. In addition, on PEZY-SC2 based systems we used single-precision calculation for the interaction kernel. In order to avoid the large roundoff at the first subtraction of the position vectors, both positions and velocities are shifted with the new origin at the position of one of the particles which share the interaction list. After this shifting, positions and velocities are converted to single precision, and actual interaction calculation is done using single-precision SIMD operations.

### IV. MEASURED PERFORMANCE

#### A. How the performance is measured

To measure the performance, we measure the time for 64 timesteps, including the time for diagnostics. The execution time is measured by the MPI wallclock timer, and operation count is from the counted number of interactions calculated. Equation 2 gives the definition of the particle-particle interaction.

$$\mathbf{F}_{ij} = \begin{cases} G \frac{m_i m_j}{r_{ij}^3} \mathbf{r}_{ij} & (r_{ij} > r_{\text{coll}}) \\ \left[ G \frac{m_i m_j}{r_{\text{coll}}^3} + \frac{m_j}{m_i + m_j} \times \left( \kappa \frac{r_{ij} - r_{\text{coll}}}{r_{ij}} + \eta \frac{\mathbf{r}_{ij} \cdot \mathbf{v}_{ij}}{r_{ij}^2} \right) \right] \mathbf{r}_{ij} & (r_{ij} \leq r_{\text{coll}}) \end{cases} \quad (2)$$

with  $\mathbf{r}_{ij} = \mathbf{r}_j - \mathbf{r}_i$ ,  $\mathbf{v}_{ij} = \mathbf{v}_j - \mathbf{v}_i$ ,  $r_{ij} = \|\mathbf{r}_{ij}\|$

Here,  $\mathbf{F}_{ij}$  is the acceleration of particle  $i$  due to particle  $j$ ,  $\mathbf{r}_{ij}$  and  $\mathbf{v}_{ij}$  are the relative position and velocity vectors,  $G$  is the gravitational constant (taken to be unity in this paper),  $m_i$  is the mass of particle  $i$ ,  $r_{\text{coll}}$  is the distance at which two particles collide, and  $\eta$  and  $\kappa$  are parameters which determine the coefficient of restitution. We chose these parameters so that the coefficient of restitution in radial direction is 0.5.

We used this form to calculate all particle-particle interaction. For particle-tree-node interaction, we used center-of-mass approximation. Particle-particle interaction consists of 9 multiplications, 8 additions, and one square root and one division operations. Instruction set of Sunway 26010 processor does not include fast approximation for neither square root or reciprocal square root. So we implemented fast initial guess and high-order convergence iteration in software. The number

TABLE I  
INITIAL CONDITION FOR WEAK SCALING RUNS

|                                 |                |
|---------------------------------|----------------|
| Central planet                  | Saturn         |
| Ring inner radius               | $10^5$ km      |
| Ring width                      | 100 km         |
| Number of MPI processes         | 1024 – 160,000 |
| Number of particles per process | $10^7$         |
| particle radius                 | 3.5 – 500 m    |

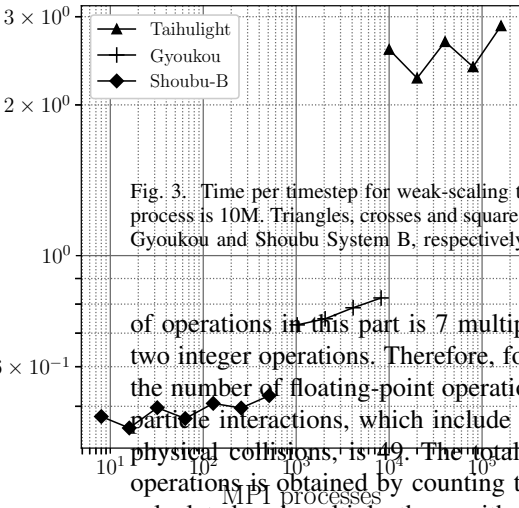


Fig. 3. Time per timestep for weak-scaling test. The number of MPI processes is 10M. Triangles, crosses and squares show the results on TaihuLight, Gyoukou and Shoubu System B, respectively.

of operations in this part is 7 multiplications, 5 additions and two integer operations. Therefore, for particle-cell interactions the number of floating-point operations is 31, and for particle-particle interactions, which include the repulsive physical collisions, is 49. The total number of operations is obtained by counting the number of operations calculated and multiply them with these number of floating-point operations per interaction. We ignore all of the operations other than the interaction calculation, since as far as the number of floating-point operations is concerned, that calculation is more than 99% of total operation.

For PEZY-SC2 based systems we used the same initial condition as we used for TaihuLight, in order to make a direct comparison possible, even though the implementation of the force kernels are different.

For the weak-scaling measurement, we have performed runs with 10M particles per MPI process on TaihuLight and PEZY-SC2 based systems. Initial condition is such that the ring radius is unchanged. Table I summarizes the initial condition.

### B. Performance Results

Figures 3 and 4 show the time per one timestep and the performance, for the weak scaling measurements. We can see that the weak scaling performance is quite good on both of TaihuLight and PEZY-SC2 based systems. The peak performance of single MPI process is roughly four times faster

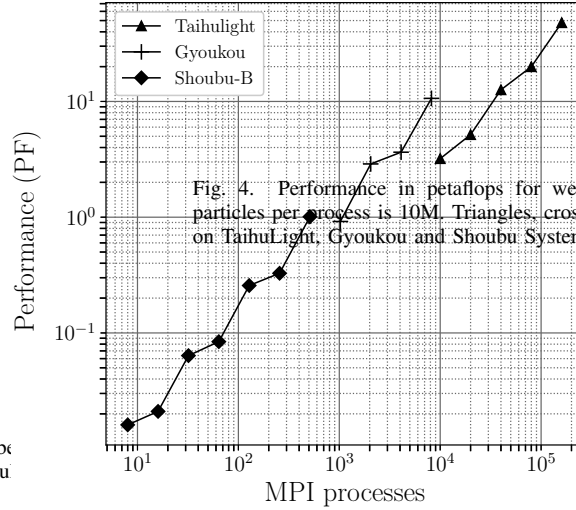


Fig. 4. Performance in petaflops for weak-scaling test. The number of particles per process is 10M. Triangles, crosses and squares show the results on TaihuLight, Gyoukou and Shoubu System B, respectively.




Fig. 5. Time per timestep for strong-scaling test. Triangles and crosses show the results on TaihuLight and PEZY-SC2 based systems, respectively. The total number of particles is  $10^{11}$  on TaihuLight and  $10^{10}$  on PEZY-SC2 based systems.

for PEZY-SC2 based systems, and that's the reason why they are three to four times faster in this weak-scaling measurement. We can see that Shoubu System B is about 50% faster than Gyoukou. As we have already discussed, this is not due to any hardware difference but the difference in the software used.

Figure 5 shows the strong scaling result. The total number of particles is  $10^{11}$  and  $10^{10}$ , for TaihuLight and Gyoukou. We do not show the strong-scaling result for Shoubu System B, since it is rather small system and strong-scaling result is not so meaningful. We can see that speedup is almost linear.

Tables II and III show the breakdown of the calculation time per one timestep, again for both the weak and strong scaling runs. As expected, in the case of strong-scaling runs, the calculation time for communication does not decrease

TABLE II  
BREAKDOWN OF CALCULATION TIME FOR WEAK-SCALING RUNS

| System     | # processes | interaction | comm. | others |
|------------|-------------|-------------|-------|--------|
| TaihuLight | 10000       | 2.07        | 0.041 | 0.466  |
|            | 20000       | 1.63        | 0.040 | 0.590  |
|            | 40000       | 2.13        | 0.064 | 0.478  |
|            | 80000       | 1.71        | 0.053 | 0.630  |
|            | 160000      | 2.31        | 0.090 | 0.476  |
| Gyoukou    | 1024        | 0.332       | 0.114 | 0.281  |
|            | 2048        | 0.392       | 0.121 | 0.235  |
|            | 4096        | 0.355       | 0.143 | 0.289  |
|            | 8192        | 0.453       | 0.147 | 0.222  |
| Shoubu B   | 8           | 0.327       | 0.018 | 0.132  |
|            | 32          | 0.344       | 0.020 | 0.181  |
|            | 128         | 0.348       | 0.027 | 0.132  |
|            | 512         | 0.360       | 0.030 | 0.135  |

TABLE III  
BREAKDOWN OF CALCULATION TIME FOR STRONG-SCALING RUNS

| System     | # processes | interaction | comm.  | others |
|------------|-------------|-------------|--------|--------|
| TaihuLight | 10000       | 2.0738      | 0.0410 | 0.4658 |
|            | 20000       | 1.0499      | 0.0253 | 0.2426 |
|            | 40000       | 0.5565      | 0.0298 | 0.1125 |
|            | 80000       | 0.2991      | 0.0233 | 0.0652 |
|            | 160000      | 0.1765      | 0.0322 | 0.0356 |
| Gyoukou    | 1024        | 0.3323      | 0.1140 | 0.2808 |
|            | 2048        | 0.1512      | 0.0668 | 0.1658 |
|            | 4096        | 0.0854      | 0.0417 | 0.0923 |
|            | 8192        | 0.0538      | 0.0357 | 0.0582 |

significantly, and eventually limits the performance. As already stated, our main interest is to use very large number of particles. Therefore, for actual scientific runs, the communication time would not become the limiting factor.

If we compare the calculation times on Gyoukou and Shoubu System B, we can see that the times for the interaction calculation are similar. but for both communications and “others”, Shoubu System B is much faster. Again, this is not due to hardware difference but due to software difference.

The performance of run for  $1.6 \times 10^{12}$  particles on 160k processes (40000 nodes) of TaihuLight is 47.9 PF, or 39.7% of the theoretical peak performance of the Sunway TaihuLight system. On PEZY-SC2 based systems, we achieved 10.6PF for  $8 \times 10^9$  particles on 8K SC2 chips, or efficiency of 23.3% of the theoretical peak performance. On 512-chip Shoubu System B, we achieved the speed of 1.01 PF, or 35.5%

The overall efficiency we achieved on PEZY-SC2 based systems is a bit lower compared to that on TaihuLight. This difference is not due to any fundamental difference in the architecture but purely due to the limitation on the available time for program development and performance measurement. As we stated, the calculation in the list-construction step, such as the constructions of the tree and the interaction lists are currently done on Xeon D, and around 40% of the total time is consumed in this step at the time of measurement on Gyoukou. Most of these are now done on SC2 side, and that is why the performance of Shoubu System B is better than that of Gyoukou.

In terms of the number of particles integrated per second, we have achieved  $5.5 \times 10^{11}$  particles per second, which is

more than 10 times faster than the results of previous works on K computer [?] or ORNL Titan [?].

## V. PERFORMANCE PORTABILITY

We have reported the measured performance of two rather different HPC systems, Sunway TaihuLight and PEZY-SC2 based systems, for the same large-scale simulation of self-gravitating planetary rings. In both cases, we have achieved fairly high efficiency, more than 30% of the theoretical peak. The parallel algorithm used is essentially the same for the two systems. However, the actual codes are rather different, simply because of the difference in the architecture and the software development environment.

Sunway TaihuLight has a heterogeneous many-core architecture integrated in one chip. Thus, the CPU (MPE in their terms) and accelerators (CPE in their terms) share the same physical memory, but CPEs lack the data cache and need to rely on DMA controller to access the main memory efficiently,

On TaihuLight, one can use OpenACC compiler. However, in order to achieve high performance, one is practically forced to use the Athread call, which makes the 64 CPEs and their local memories visible to programmers.

On the other hand, PEZY-SC2 systems, at least at present, have a rather classical accelerator-based architecture, in which CPU (a Xeon-D processor) and accelerators (PEZY-SC2 processors) are connected through PCI Express interface. This means that they have separate physical memories. Within one chip, however, processing elements of PEZY-SC2 processor have three levels of data caches. Currently PZCL, a dialect of OpenCL, is supported on PEZY-SC2 based systems.

Because of these differences (shared and separate memory, DMA and cache, thread-based and OpenCL-like), the actual programs for two machines have become quite different, even though the algorithms used are the same and the problem to be solved is the same.

Both codes, however, are based on our framework, FDPS [?] and follow its structure. The basic idea of FDPS is to separate the implementation of parallel algorithms and description of the physical problem. FDPS provides the former and the application programmers provides the latter, in the form of the data type definition of particles and functional form of particle-particle interaction. Users of FDPS can write their programs by specifying the data structure of particles they use, and calling necessary FDPS functions for domain decomposition, particle migration between processes, and interaction calculation. currently, users should provide optimized function for particle-particle interaction calculation.

Many of the parallel algorithms we newly implemented are not specific to planetary rings but can be applied to any other particle-based simulations. Using FDPS, users can write their programs in their favorite language (currently C++ and Fortran are supported and we plan to extend the language set), and let FDPS do complex parallelization.

Thus, it seems that one way to achieve performance and program portability on new machines with rather exotic architecture such as the machines evaluated in this paper is

to develop the framework with a common API and internal implementations specialized and optimized to specific architectures. This is fairly straightforward in the case of TaihuLight, in which the CPE and MPEs share the single physical memory, since the data structure that FDPS handles can still in the shared main memory. The basic data structure of FDPS is just an array of particles, and both the user-developed application program and the FDPS side can access that particle array in usual way.

On the other hand, how the separate memory spaces of PEZY-SC2 should be handled within FDPS requires a bit more consideration. One possibility would be to add the interface in which the user-side programs, for example the function to perform time integration, is passed to FDPS, instead of directly called within the user-side program. Here, the function to be passed applies to single particle (or some small array or particles), and applying it to all particles in the system will be the responsibility of FDPS. This approach will probably make the software development and performance improvement easier on other machines, since parallelization in both MPI and OpenMP level can be taken care of within FDPS.

This problem of portability is of course not limited to FDPS. It occurs in practically any application in any field of computational science. We clearly need a new and systematic approach to solve this problem, and we think the use of frameworks such as FDPS may be an efficient and practical way.

Our view of frameworks is that they should allow users to express their problems in simple and (ideally) machine-independent way. In the case of particle-based simulations, We have designed FDPS to meet this goal, and it actually works pretty well on large HPC systems, both with and without GPGPUs. Our experience on TaihuLight and PEZY-SC2 indicates that it is also possible to extend FDPS to cover these systems. We believe similar approaches will be used in other fields.

## VI. SUMMARY

In this paper, we described the implementation and performance of a highly efficient simulation code for self-gravitating planetary rings on Sunway TaihuLight and PEZY-SC2 based systems.

The measured performance is 47.9 PF, or 39.7% of the theoretical peak, for simulation of  $1.6 \times 10^{12}$  particles on 40,000 nodes of TaihuLight, 10.6PF, or 23.3% of the theoretical peak, for simulation of  $8 \times 10^{10}$  particles on 8192 nodes of Gyoukou, and 1.01PF, or 35.5% of the theoretical peak for  $5 \times 10^9$  particles on 512 nodes of Shoubu System B. As noted earlier, Gyoukou and Shoubu System B use the same PEZY-SC2 processor. The difference in the efficiency is purely due to the fact that Gyoukou was turned off on March 31, 2018. The software at that time was still under development.

Compared to previous achievements on K computer or ORNL Titan, the achieved efficiency is similar or higher, and the speed in terms of the number of particles integrated per second is higher, for both TaihuLight and PEZY-SC2 based

systems. As we stated earlier, this level of performance would not be achieved without the new algorithms described in this paper.

Compared to other multi-core processors for modern HPC systems such as Fujitsu SPARC64 VIIIfx and IXfx or Intel Xeon Phi processors, both SW26010 processor of TaihuLight and PEZY-SC2 processor of PEZY-SC2 based systems have several unique features which allow very high peak performance but at the same time make it much harder to achieve high efficiency on real applications. These are:

- Heterogeneous architecture with rather extreme performance ratio of 1:64 in the case of SW26010 and even larger in the case of SC2.
- The lack of cache hierarchy (SW26010) or cache coherency (SC2).
- Very limited main memory bandwidth, with B/F values around 0.02–0.04. This is about 1/10 of the numbers of Fujitsu or Intel HPC processors.

On the other hand, SW26010 comes with very well-thought features which allows the programmers to optimize the performance of code on CPE. These features include:

- Low-latency DMA controller which can be initiated by any CPE.
- Low-latency, high-bandwidth communication between CPEs.

These two features allow very efficient use of the main memory bandwidth. The two-dimensional structure of the network within CG seem to be optimized for highly efficient implementation of matrix-matrix multiplications, but it is actually quite useful for other real applications, whenever fast inter-core communication is necessary.

It is certainly true that the need to use DMAs for data transfer between CPE and main memory complicates the use of CPE. However, it is also true that it makes quite optimized access to main memory possible, since the application programmer can (or have to) control all main memory accesses. In the case of our code, in several places we have “vectorizable” loops, which perform the same operation on all particles in the system. The number of operations per particle is relatively small, of the order of ten, and the data size of one particle is 32 bytes. In the case of manycore architecture with hierarchical cache memory, to achieve high efficiency on simple vector operations like

```
a[i] = b[i]+c[i]
```

is actually quite complicated. In modern processors, load address would be predicted and hardware prefetch is generated. The hardware prefetch would probably work for a very simple loop like the above example, but would fail if many vectors are loaded. Then the programmer need to experiment with software prefetch, to find the way to get the best performance.

In the case of SW26010, currently it is rather tedious and error-prone to write the equivalent operation using the combination of Athread and DMA, and sometimes inner kernel in assembly language, but once we do so, we can get a performance close to the theoretical limit relatively easily.



The existence of low-latency (less than 10 clock cycles) communication path between CPEs is quite important for using CPEs for fine-grain parallelism such as loop-level parallelization. Such low-latency communication is difficult to implement on shared memory processors with hierarchical cache.

The SC2 processor supports the cache flush and synchronization at each level of the cache hierarchy, making the relatively low-latency communications between processors possible. However, it is clearly desirable to have more direct control of interprocessor communication.

One common problem of SW26010 or SC2 is that writing high-performance kernel for them means writing the inner kernel in the assembly language. This is purely the software limitation, and probably not so difficult to fix. In this aspect, SC2 is somewhat easier to deal with, since it supports 8-way multithreaded execution, which effectively hides the latencies of L1 cache and arithmetic unit from compiler.

In conclusion, we have implemented parallel particle simulation code on Sunway SW26010 and PEZY-SC2 processors, and found that it is not impossible to achieve high performance on their rather extreme architectures with carefully designed algorithms. Even though the B/F number are less than 0.1 and the network bandwidth is similarly low, the efficiency we have achieved is comparable to that on K computer, with 15 times more memory and network bandwidth. We feel that architecture evolution in this direction will help the HPC community to continue improving the performance. We also believe that high-level software framework such as our FDPS will help many researchers to run their own applications efficiently on new architectures.

#### ACKNOWLEDGMENT

This work was supported in part by The Large Scale Computational Sciences with Heterogeneous Many-Core Computers in grant-in-aid for High Performance Computing with General Purpose Computers in MEXT (Ministry of Education, Culture, Sports, Science and Technology-Japan) and also in part by JAMSTEC, RIKEN, and PEZY Computing, K.K./ExaScaler Inc. In this research computational resources of the PEZY-SC2 based systems supercomputer, developed under the Large-scale energy-efficient supercomputer with inductive coupling DRAM interface project of NexTEP program of Japan Science and Technology Agency, has been used.